

BSc. COMPUTER SCIENCE



UNIVERSITY OF AMSTERDAM

Modernizing the Civic compiler framework

Floris den Heijer
5873355

June 12, 2015

Supervisor(s): dr. C. Grelck

Signed:

Abstract

Abstract, todo

Contents

1	Introduction	5
1.1	Problem Statement	5
1.2	Context	5
1.3	Objectives	5
1.4	Research Approach	5
1.5	Scope And Limitations	5
2	Analysis Of Civic Framework	7
2.1	Architecture	7
2.2	Key Concepts	8
2.2.1	Seperate Specification And Implementation	8
2.2.2	Flexible Traversals	9
2.2.3	Tree Validation	9
2.3	Development Environment	9
2.4	Strengths & Weaknesses	9
3	Design Considerations	11
3.1	Language	11
3.2	Architecture	12
3.3	Code-First vs. Specification	12
3.4	Validation	12
3.5	Traversals	12
3.6	Collections	12
4	Framework Changes	13
4.1	Code Defined AST	13
4.2	Pluggable Phases and Traversals	13
4.3	Improved Traversal Mechanism	13
4.3.1	Read-Only Visitor	14
4.3.2	Extrinsic Visitor	14
4.3.3	Lambda Visitor	16
4.4	Development Environment	17
4.4.1	Updated Lexer/Parser	17
4.4.2	Platform Updates	18
4.4.3	Documentation Generator	18
5	Discussion	19
6	Bibliography	21

Introduction

1.1 Problem Statement

Its so hard and confuzing!! [1]

1.2 Context

Purpose of framework, teaching objectives, timeline.

Other compiler frameworks, LVVM as backend, research teaching methods for compiler construction??

1.3 Objectives

The primary goal of this research project is to investigate how the Civic framework can be simplified by transitioning to a more modern language.

Key questions answered:

- Which aspects of the current framework arise from choice of language, and which from methodology?
- What are the drawbacks and advantages of the chosen platform and architecture?
- If a change of language is permitted, how can the framework be improved?

1.4 Research Approach

Analysis of existing framework, consideration of language dependent and independent factors, limitations of analysis.

1.5 Scope And Limitations

The Civic framework is used in a relatively well-defined educational setting which omits many aspects of more advanced, general purpose compiler frameworks. Teaching objectives dictate the primary focus to be on the development of a functional frontend (scanning, parsing, semantic analysis), rather than a complete backend. Focus is placed on creating a flexible abstract syntax tree which is suitable for basic optimization and code emission, eliminating the need for a separate intermediate representation. While greatly simplifying the compiler pipeline, this reduces the framework to a strictly 'frontend-only' role and limits comparisons with other frameworks.

This paper is limited to the Civic-to-Civic VM portion of the framework and will not cover analysis of the Civic VM, assembler, teaching objectives nor the suitability of alternative target languages such as LVVM, MSIL or Java bytecode.

Analysis Of Civic Framework

This section will provide an overview of the architecture of the existing framework and its relevant components. On a high level, the framework aims to provide a structured approach for phase-based transformations on a language agnostic AST definition. Key components of the framework are:

1. Abstract Syntax Tree definition
2. Code generation based for node access and traversals
3. Transformation pipeline

In addition it provides a documentation generator for the AST, a small base library for string manipulation, lookup tables and argument parsing and a testing suite. The framework is written in C and XLST and runs on most *nix distributions, provided they support the required packages¹.

2.1 Architecture

The outline of the architecture is presented in figure 2.1. It makes heavy use of XSLT to transform abstract specifications for the AST, traversals and node composition into usable helper code. This code combined with the base library provides the macro's and header files required for the rest of the compiler. The Flex/Bison lexer/parser is automatically compiled and transformed into a usable phase. Next, user phases and additional code are compiled and traversal entry points are substituted in the phase specification file and included in the main compilation unit.

¹gcc, gzip, flex, bison, xsltproc, dot, indent

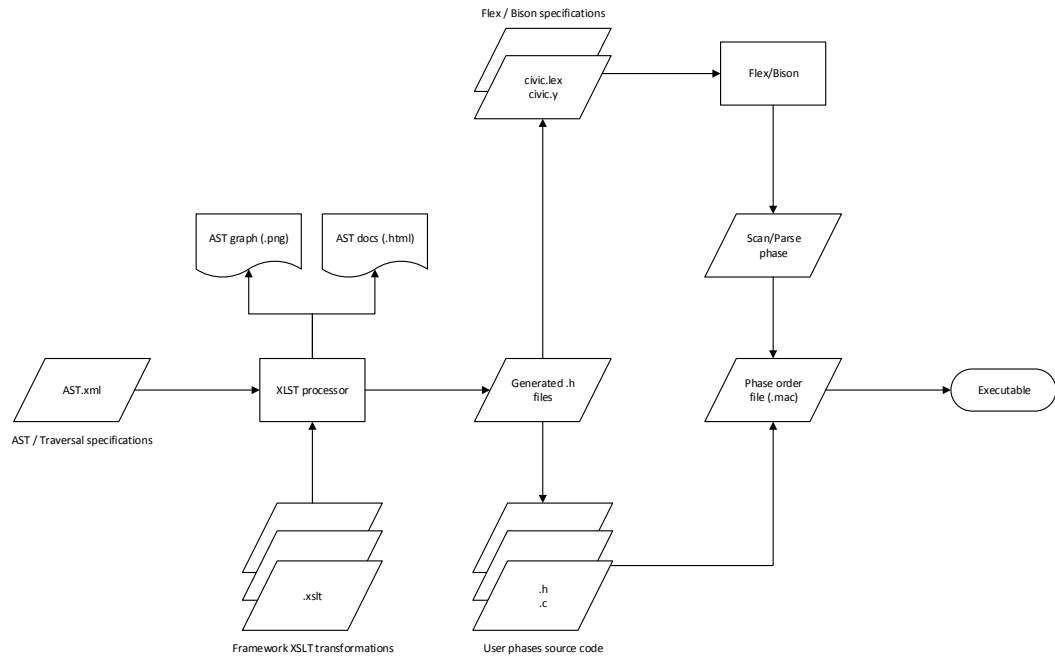


Figure 2.1 *Simplified overview of the Civic framework*

2.2 Key Concepts

2.2.1 Separate Specification And Implementation

Central in the approach taken by the framework is the separation of specification and implementation. Specifications are given for:

- Nodes which form the abstract syntax tree, their attributes and child nodes
- Mapping from friendly attribute names to C-types, and relevant methods used for copying and free-ing
- Node constraints on:
 - Type: only certain child nodes or attribute are valid
 - Requirement: attributes or child node must be present
- Traversals, their mode of operation, affected nodes and function prefix

These specifications are transformed with XSLT to:

- Access macro's for nodes and their attributes
- Methods for initializing nodes
- Validity assertions
- User defined tree traversals and system traversals:
 - Check traversal
 - Free traversal
 - Copy traversal
- AST documentation (HTML) and graph (PNG)

2.2.2 Flexible Traversals

How is a tree represented (internally), and how to traverse?

2.2.3 Tree Validation

How is a tree validated, and how are type errors found?

2.3 Development Environment

Error handling, pre-configured lexer+parser, debug helpers, utilitarian base library

2.4 Strengths & Weaknesses

Weaknesses:

- C language probable cause for:
 - Complicated traversal notion (due to macro's)
 - Check traversal could be made obsolete by including type checks in macro's
 - Free traversal (see GC)
 - Mixing of C-related implementation details in specification
 - Frequent use of debug helpers (doesn't help RAD)
 - Ancient debugging, verbose code, everything global, null pointers, etc.
- Platform probable cause for:
 - Extensibility issues (XSLT)
 - Architecture not obvious
 - Large specification required for small compiler (huge XML)
 - Documentation (due to perceived complexity)

Strong points:

- Everybody knows C (but nobody knows the framework)
- Platform is widely supported
- Forces rigid structure, deviation impossible due to non-transparency

Design Considerations

Modern language, infrastructure approach, type-safety, code-based AST. Key questions:

- Which programming languages are realistic alternatives to the current language, keeping in mind the current experience of students?
- Which areas or concepts could be improved or introduced to improve the framework's efficiency?
 - In what way are these concepts dependent on implementation language or on choice of architecture?
 - Which concepts are invariant of platform or architecture?

3.1 Language

As established in section 2.4 a large number of shortcomings can be attributed to the C programming language.

Which programming languages are realistic alternatives to the current language, keeping in mind the current experience of students?

Options: Java, Scala, C#, C++, Python, Functional (F#, OCaml)

Considerations: Student background (no functional, Java experience), New technologies accessible (ANTLR4, advanced debugging, exception handling), ease of use

Choice: C# (ANTLR, similarity to Java with more (functional) features, availability, cross-platform footnote, personal experience, managed language, extension methods)

Implications for platforms: Windows, Mono, base-CLR, etc.

Code based configuration of phases and AST are major differences, overall similar approach. New error handling helper methods and type-safety.

Something about a modern platform.

- Change to a managed OOP language relieves memory allocation issues, simplifies code based AST definition
 - Direct control over AST
 - Inherent type safety
 - More flexibility in choice of parser/lexer generator
 - Base library removes need for wrapper methods and guarantees documentation
- Code configured pipeline allows easier manipulation of phases
- Action based traversal formalizes and extends existing traversal process with conditionals and inheritance based actions

3.2 Architecture

Phase based is common and works well in general (proof?), AST graph makes sense (as does docu),

3.3 Code-First vs. Specification

Main motivator...? Disadvantages and limitations (better suited for large-scale? multiple ASTs?)

The current framework fails to completely separate specification from implementation. But even an imaginary framework which succeeds has to interface with code at some point.

Considering code and AST are directly related, improvement is to construct AST from C# classes, offering type safety across traversals, inheritance options and more extensibility when compared to XSLT generated code.

Honorable mentions: interface based approach, attribute tagging for order of child nodes, code generation based on inheritance graph (for visitors)

Language dependency: Any managed OOP language supports this construct, though C# additionally allows easy construction of object graph using bracket notion.

Platform dependency: Breaks with AST XML definition.

3.4 Validation

Type-safety is more than enough?

3.5 Traversals

Optimal way to traverse and modify a tree? Visitor pattern, algorithm separate from structure, not much choice.

Outcome: stuck with some form of visitor, more in section x..

3.6 Collections

When given a strongly typed object graph, why not utilize the built-in support of list like structures? Discussion of problems.

Case with collections: - Node child may be node (or derived) or list of nodes - Discriminated union in C# (<https://stackoverflow.com/questions/3151702/discriminated-union-in-c-sharp>) - What doesn't work: Every node is also a list of nodes - Cyclic dependency in inheritance (Node : IList?) - Advantages: easy to add, iterate and replace - Disadvantages: Traversal issues (can a handler return a list?), complicates everything

Case without collections: - More natural to traversals - Singly-linked or doubly-linked - Difficult to append to or modify - Can be overcome with extension methods

Outcome: chosen for no collection types in lists, to avoid confusion with traversals.

Framework Changes

Fundamentally changed:

- Code-First vs specification (for AST, traversals and phases - also validation omitted)
- Managed Language
- Package updates

Untouched or adapted:

- Traversal Mechanism (added lambda)
- Collection representation
- Base lib (arguments, default ANTLR4, documentation)

4.1 Code Defined AST

...

4.2 Pluggable Phases and Traversals

Pluggable traversals, not the actual rework. Pluggable phases, interface definition, etc.

4.3 Improved Traversal Mechanism

AST transformation through tree traversals were and still are key components of a compiler framework. The original framework relied on configuration based code generation to create the necessary traversal paths and helper methods, resulting in indirect and cumbersome configuration of traversals.

AST transformation is crucial and diverse. Operations such as type checking, code generation, variable initialization checks, etc. Could be implemented as *TypeCheck()* on objects, but this doesn't scale. Type checking operations should be kept together, etc.

Extrinsic visitor: Easy addition of new methods, separation of dissimilar operators, concrete elements visited is small, trade runtime performance for reduced intrusiveness of visitor pattern (Nordman)

Walkabout is exactly what I built (essence of visitor pattern).

When transforming an AST, it makes sense to use a type of depth-first traversal as actions on child nodes are often dependent on the root node and less on their cousins.

The visitor pattern is a well established software pattern which separates algorithm from object structure. Specifically, it tackles traversal of a structure whose objects are part of an inheritance hierarchy.

In the traditional double-dispatch version this is done by adding an *accept(IVisitor)* method to all participating objects which directly return control to the visitor, invoking the *visit(Object)* method on the visitor determined by the run-time type (Essense of the visitor pattern).

There are many variations on the visitor pattern, such as the extrinsic visitor pattern (Norman), reflection, type test and delegate (C# patterns), hierarchical etc (C2.com).

Introduction to composite and visitor pattern, example with no inheritance.

Review of existing hybrid visitor pattern, does not translate to with inheritance. Visitor may return more generic node than given. Invoked method depends on declared node type (single dispatch).

Double dispatch, calling visitor method with runtime type. However: - does not allow handling of types higher up in the inheritance chain - not natively supported by C#, even possible?

Alternatives (unsure): - store parent in every node, but who is responsible for setting it? Lifts responsibility to handler. Which property handled?

Introduce case with multiple matching methods, implies declaration order matters.

Define handler: - type in - type out - action or function (can be method invocation)

Define visitor: - ordered list of handlers - generic visitor method with reflection lookup

Added features: - type out omitted -> wrap with return of same type - conditionals on handler declaration

AbstractA / ConcreteB ConcreteC | ConcreteB2

Visitor with: AbstractA Visit(AbstractA node) ConcreteB Visit(ConcreteB node) AbstractA Visit(ConcreteB2 node)

Goal: preserve type safety in visitor methods without using casts: wrong: `expr.Left = (Expression)Visit(expr.Left)`; good: `expr.Left = Visit(expr.Left)`;

4.3.1 Read-Only Visitor

...

4.3.2 Extrinsic Visitor

Suppose transformation of IntConst to Cast (both Expression) and default visit method: `Expression Visit(IntConst node)`; `T Visit(T node)`;

Assume we have a binop node with properties Left and Right of type Expression:

Snippet 4.1 *Spanning tree broadcast algorithm.*

```
1 var int1 = new IntConst();
2 var int2 = new IntConst();
3 var binop = new BinOp { Left = int1, Right = int2 };
```

Invoking the visitor like this results in invocation of the default Visit handler.

Snippet 4.2 *Spanning tree broadcast algorithm.*

```
1 binop.Left = Visit(binop.Left);
```

Lets follow the Walkabout pattern for the generic visitor method:

Snippet 4.3 *Spanning tree broadcast algorithm.*

```
1 T Visit(T node) where T : Node
2 {
3     if (node != null)
4     {
5         // If this visitor has a public method for this class, with convertible return
           type, invoke.
```



```

6     var nodeType = node.GetType();
7     var method = this.GetType().GetMethod("Visit", new[] { nodeType });
8
9     if (method != null && (method.ReturnType == typeof(T) ||
10        method.ReturnType.IsSubclassOf(typeof(T))))
11     {
12         return (T) method.Invoke(this, new [] {node});
13     }
14     // Otherwise traverse children.
15     VisitChildren(node);
16 }
17
18 return node;
19 }

```

Suppose default visit children method which traverses and sets child node properties:

Snippet 4.4 *Spanning tree broadcast algorithm.*

```

1 void VisitChildren(Expression expr)
2 {
3     foreach (prop of expr.GetNodeDerivativeProps())
4     {
5         Node value = (Node) prop.GetValue();
6         if (value != null)
7         {
8             value = Visit(value);
9             prop.SetValue(newValue);
10        }
11    }
12 }

```

GetValue() and SetValue operate on objects, the cast to Node is only required to ensure the most general case of T Visit(T node) functions.

Invoking the VisitChildren method on the BinOp successfully replaces both IntConst's with Casts.

Now suppose we want to write a method which acts on ALL expressions: Expression Visit(Expression expr);

We need to establish rules on precedence. The most intuitive approach would be to assume the most concrete visitor method should be called, similar to single-dispatch.

Now that we have handlers on two inheritance levels, invoking the most concrete handler involves a dynamic cast.

Snippet 4.5 *Spanning tree broadcast algorithm.*

```

1 binop.Left = Visit((dynamic) binop.Left);
2 binop.Right = Visit(binop.Right);

```

VisitChildren also works because the GetMethod reflection call takes into account inheritance hierarchy.

But is this approach type safe? A dynamic object is essentially something of type 'object' and exists only at compile time (MSDN). The fundamental weakness in this approach lies in the premise that a node handler can return something more abstract than itself. To introduce a run-time error all we have to do is make the IntConst handler more abstract by returning a node: Node Visit(IntConst node) return new Node();

This runtime error is thrown because the property on BinOp requires an Expression and the generic visit method has not taken this into account. The value of a property is always returned as an object, and the property field type is retrieved using reflection. This means we cannot cast to the required type resulting in a perfectly working visitor returning an object too generic.

It is possible to supply this information to the Visit method without exposing it to outer classes. First the generic visitor and VisitChildren methods will be placed in an abstract class

from which a new concrete visitor will inherit. Next, the visitor function is modified and made private:

Snippet 4.6 *Spanning tree broadcast algorithm.*

```
1 T Visit(T node, Type maxUpcast) where T : Node
2 {
3     ...
4     if (method != null && (method.ReturnType == maxUpcast ||
5         method.ReturnType.IsSubclassOf(maxUpcast)))
6 }
```

A wrapper method is introduced which hides this and restores the original behavior:

Snippet 4.7 *Spanning tree broadcast algorithm.*

```
1 T Visit(T node) where T : Node
2 {
3     return Visit(node, typeof(T));
4 }
```

Finally the VisitChildren function is adjusted to pass along the type of the property:

Snippet 4.8 *Spanning tree broadcast algorithm.*

```
1 value = Visit(value, prop.PropertyType);
```

Instead of throwing a runtime error, the visitor now assumes the default case, skips invocation and jumps straight to visiting children. Static type checking is still performed by the C# compiler, preventing this from compiling:

Snippet 4.9 *Spanning tree broadcast algorithm.*

```
1 binop.Right = visitor.Visit((BoolConst)binop.Right);
```

The only case which is not supported in this scenario is one where there are two matching methods with differing return types:

Snippet 4.10 *Spanning tree broadcast algorithm.*

```
1 Expression Visit(IntConst i);
2 IntConst Visit(IntConst i);
```

Drawbacks: always have to invoke visitor with dynamic cast when also targeting higher up the inheritance hierarchy.

4.3.3 Lambda Visitor

While this approach is sufficient for most traversals, it requires some care when multiple inheritance layers are targeted. Method invocation depends on two separate mechanisms, namely direct (single-dispatch via compiler) and reflection based. This ambiguity can be avoided by preventing direct access to visitor methods. This approach takes the decoupling from the extrinsic visitor design one step further by removing publicly defined visitor methods and replacing them with configuration functions which attach handlers.

The basic traversal mechanism is the same, except instead of probing for publicly exposed methods an ordered list of handlers is consulted instead. This also prevents ambiguity when multiple layers of inheritance are targeted by explicitly stating the order of evaluation. Such handlers must contain at least:

1. Type of the node handled
2. Return type of the handler (similar to the upcast restraint)
3. Delegate containing the transformation function

The previous distinction between methods which replace nodes and those who only act on them can be removed by offering two types of delegates. In .NET these are represented by the `Func<T, T>` and `Action<T>` delegates, indicating respectively a function which takes an object of type `T` and returns an object of type `T` and a function which takes an object of type `T` but doesn't return anything. These can be interleaved as in returning nothing in essence indicates no change to the initial object; it would behave 'as-if' it would return itself.

Considering the logic determining which visitor function should be invoked is now entirely encapsulated by the visitor itself, a simple but useful feature would be an optional predicate preceding invocation of the function. This, together with a lambda-style approach to defining visitor methods allows for more configuration flexibility. Such handlers must additionally hold:

4. Return type of the delegate
5. Delegate containing a predicate

Snippet 4.11 *Test*

```

1  @misc{website:sac,
2    author = "SAC-Home",
3    title = "Homepage of the SAC project",
4    month = "May",
5    year = "2015",
6    url = "http://www.sac-home.org/"
7  },
8  @article{unknown,
9    author = "Source?",
10   title = "Everything Explained Vol. 2",
11   month = "March",
12   year = "2012"
13 }
```

Additionally, fall-through handlers could be supported with this approach. The decision not to do so is based on the fact that fall-through statements have been mostly banned from the switch construct in modern imperative languages such as Java and C#. The recent HeartBleed-bug in OpenSSL was also caused by an improperly handled fall-through.

Suppose we have two handlers (in-order): `F(IntConst, Node)` `F(IntConst, Expression)`

In the `BinOp` example previously defined the second handler would be invoked as they are held in properties of type `Expression`.

Adding

4.4 Development Environment

Implications of unmanaged vs managed, reduction of base library, debug upgrade..?

4.4.1 Updated Lexer/Parser

More flexible Parser/Lexer generator support, such as ANTLR4 (LL(*) with EBNF).

Language dependency: Plenty of parsers exist for any language, first version of AST most likely constructed in code; change of language not essential.

Platform dependency: Depends on extensibility of parser generator, ANTLR4 has multiple ways of traversing a generated parse tree, YACC does not (embedding of code in parse phase restricts construction of AST)

4.4.2 Platform Updates

NuGet

4.4.3 Documentation Generator

Inheritance graph plotter, AST graph plotter

Discussion

Key questions:

- Which aspects of the existing framework are a result of language?
 - How does the new framework differ from other compiler platforms?
 - What other factors could improve the efficiency of the framework?
- vNext immutable AST -

CHAPTER 6

Bibliography

[1] Source? Everything explained vol. 2. March 2012.